

Gravity works on big scales

INTRINSIC TIDAL SHEAR AT THE LARGEST SCALES: GALAXY MULTIPLY ALIGNMENT IN DESI DR2

CLAIRE LAMMAN^{1,2,3}, JESSICA NICOLE AGUILAR⁵, STEVEN AHLEN⁶, ALEJANDRO AVILES^{7,8}, DAVID BROOKS⁹, AURELIO CARNERO
ROSELL^{10,11}, TODD CLAYBAUGH⁵, AXEL DE LA MACORRA¹², DANIEL EISENSTEIN¹³, JAIME E. FORERO-ROMERO^{36,37}, CRISTHIAN
GARCIA-QUINTERO¹³, ENRIQUE GAZTAÑAGA^{14,15,16}, GASTON GUTIERREZ¹⁷, KLAUS HONSCHIED^{1,3,4}, MUSTAPHA ISHAK¹⁸, TANVEER
KARIM¹⁹, DAVID KIRKBY²⁰, OFER LAHAV⁹, MARTIN LANDRIAU⁵, MARC MANERA^{21,22}, AARON MEISNER²³, RAMON MIQUEL^{24,22},
HERNAN ENRIQUE NORIEGA^{8,12}, WILL PERCIVAL^{25,26,27}, MIGUEL PERDOMO³⁶, FRANCISCO PRADA²⁸, IGNASI PÉREZ-RÀFOLS²⁹,
CORENTIN RAVOUX³⁰, ASHLEY J. ROSS^{1,2,4}, GRAZIANO ROSSI³¹, ROSSANA RUGGERI³², EUSEBIO SANCHEZ³³, CHRISTOPH SAULDER³⁴,
DAVID SCHLEGEL⁵, JOSEPH HARRY SILBER⁵, MAŁGORZATA SIUDEK^{16,11}, GREGORY TARLÉ³⁵, BENJAMIN ALAN WEAVER²³

Affiliations are listed at the end of the paper

Version September 3, 2026

ABSTRACT

We measure the intrinsic alignment (IA) of galaxy multiplets with the large-scale tidal field traced by galaxy positions in DESI Data Release 2. Using the BGS, LRG, and ELG samples spanning $0.01 < z < 1.6$, we optimize the scales which define multiplets in each sample to minimize measurement noise. We measure multiplet tidal alignment $\mathcal{E}_+$ on scales beyond $100 h^{-1}\text{Mpc}$ in all samples, and detect alignment over $200 h^{-1}\text{Mpc}$, up to 4.5σ in the highest LRG redshift bin. We additionally present the multiplet shape–shape autocorrelation, $\mathcal{E}_{++}$, which is independent of galaxy shape, and is consistent with the nonlinear alignment model. We do not detect a baryon acoustic feature in the projected estimator. We also find the alignment strength of multiplets is correlated with redshift and galaxy morphology. Multiplet alignment offers direct access to the tidal shear field in all galaxy samples, even where individual galaxy alignment cannot, and extends IA detection to the largest scales yet probed.

1. INTRODUCTION

While often traced directly with galaxy positions, the large-scale structure (LSS) of the universe also generates a measurable gravitational field. This tidal field encodes information that is complementary to, and in some cases inaccessible from, two-point clustering. Such information is particularly valuable when measured on large scales, where the field is close to linear and retains a clean connection to primordial physics and the initial conditions of the universe. The large-scale tidal field can be used to constrain primordial non-Gaussianity (PNG), cosmic B-modes, and parity violation, particularly around scales of $100\text{--}1000 h^{-1}\text{Mpc}$, where these signatures are concentrated (Akitsu et al. 2021; Kurita & Takada 2023; Achúcarro et al. 2022; Paul et al. 2024). However, largest current direct detections of the tidal field only extend to around $100 h^{-1}\text{Mpc}$ (Figure 9).

Cosmic shear surveys probe the tidal field of LSS through the integrated effect that gravitational lensing has on background galaxies (The LSST Dark Energy Science Collaboration et al. 2018). More directly, the non-projected tidal shear field can be traced by objects that are intrinsically aligned with it, such as galaxies (Joachimi et al. 2015; Troxel & Ishak 2015; Lamman et al. 2024a; Chisari 2025). Measuring the intrinsic alignment (IA) of galaxy shapes to the underlying density is essentially correlating galaxy positions with a spin-2 object and is comparable to the collapsed form of the 3-point correlation function. This collapsed form is expected to have unique sensitivity to several of the cosmological effects listed above (Chisari & Dvorkin 2013; Kurita et al. 2025). However, measuring IA is challenging in practice; it is significantly susceptible to imaging systematics and restricted to the populations which show measurable alignment – large, red galaxies (Samuroff et al. 2023; Siegel et al. 2025; Hervas Peters et al. 2025). With cur-

rent surveys, this also functionally limits detection to below redshift $z = 1$.

Lamman et al. (2024b) presented an alternative that circumvents these issues: multiplet intrinsic alignment (MIA). This correlates the projected orientations of small sets of galaxies with the surrounding large-scale matter, as opposed to the orientations of individual ones. Galaxy multiplets typically consist of 2–5 galaxies within a few $h^{-1}\text{Mpc}$ of each other. They are not necessarily clusters, groups, or a proxy for halos. We show in this work that they often contain fewer galaxies spread over a larger volume, and may be better understood as tracing filaments. Like IA, their shapes are reduced to an orientation at a single point. It is a form of a projected, collapsed 3-point correlation function, where the collapsed side is weighted over the positions of the multiplet’s member galaxies. This configuration is motivated by our understanding of galaxy formation, where matter falls in along tidally-aligned filaments, and is considerably simpler to measure and interpret than the full bispectrum. Above $\sim 20 h^{-1}\text{Mpc}$, it can be modeled with existing IA frameworks such as the nonlinear alignment model, or NLA (Bridle & King 2007).

Using spectroscopic data from the Dark Energy Spectroscopic Instrument (DESI)’s first data release, multiplet alignment has been detected out to projected separations of $100 h^{-1}\text{Mpc}$, in populations of blue galaxies, and beyond redshift $z = 1$ (Lamman et al. 2024b). It also shows promise for cross-survey analysis: multiplets identified in imaging surveys around spectroscopic tracers display similar signals with improved signal-to-noise ratios (Kumwembe et al. 2025).

Here we explore different definitions of multiplets, identifying the scales which contain the most information and how best to weight the relative positions of member galaxies. Combined with updated data from DESI’s second data release, we extend the scales at which intrinsic tidal shear can be de-

tected and take a preliminary look at a potential BAO feature in the projected tidal shear correlation - which has yet to be detected. We also, for the first time, measure the multiplet-multiplet autocorrelation. This is a form of the collapsed 4-point correlation function and of interest for several reasons. The autocorrelation is independent of galaxy bias and provides a valuable check for systematics and the internal consistency of our model. It is also the lowest-order estimator sensitive to cosmic B-modes and certain classes of parity violation. These measurements also provide the parameters necessary to extend existing IA forecasts to multiplet alignment, including characteristic number densities, alignment strengths, and redshift dependence.

Section 2 describes the DESI DR2 catalogs. Section 3 introduces multiplets, how they are defined, and their demographics across the DESI catalogs. Section 4 summarizes the formalism and modeling of multiplet alignment. Section 5 documents our methods of measurement and error estimation, which are also performed on mock data, described in Section 6. Our main results are presented in Section 7, including large-scale measurements across all samples, autocorrelation functions, correlations of alignment strength with redshift, and investigations of a potential BAO feature.

Throughout this paper, we assume the cosmological parameters of $H_0 = 69.6 \text{ km s}^{-1} \text{ Mpc}^{-1}$, $\Omega_{m,0} = 0.286$, $\Omega_{\Lambda,0} = 0.714$.

2. DESI CATALOGS

DESI (DESI Collaboration et al. 2016, 2022) is a robotic, fiber-fed spectroscopic surveyor mounted on the Mayall 4-meter telescope at Kitt Peak National Observatory. It obtains simultaneous spectra of nearly 5000 targets across a $\sim 3^\circ$ -diameter field of view via a robotic focal plane (Silber et al. 2023), a corrected wide-field optical assembly (Miller et al. 2024), and a fiber system feeding ten three-arm spectrographs (Poppett et al. 2024). Targets are selected from DR9 of the Legacy Imaging Surveys (Dey et al. 2019; Myers et al. 2023), and observations are planned and executed by the survey operations pipeline described in Schlafly et al. (2023). Spectra are reduced and redshifts extracted with the DESI spectroscopic pipeline (Guy et al. 2023). The final survey will consist of over 63 million spectra of galaxies and quasars taken over eight years.

In this work, we use 19 million spectroscopic redshifts from DESI’s first three years of observations, Data Release 2 (DR2). The DR1 dataset and its data model are described in DESI Collaboration et al. (2026), and the DR2 samples and their cosmological analysis are presented in DESI Collaboration et al. (2025). We use the large-scale structure (LSS) catalogs (Ross et al. 2025) for all analysis. The catalogs contain weights to account for redshift failures and targeting incompleteness, which are used to weight the tracer sample in each cross-correlation measurement, and associated random catalogs matched to the survey geometry, which are also used in Section 5.

We use three LSS catalogs. The first is the Bright Galaxy Survey (BGS), which is a relatively dense sample containing galaxies with a variety of morphologies and nearly 100% completeness at low redshifts (Hahn et al. 2023). We split this catalog into five redshift samples between $0.01 < z < 0.5$. We use the full BGS catalog (“BGS BRIGHT”) rather than the absolute-magnitude limited subsample adopted for the BAO analysis (DESI Collaboration et al. 2025), which includes more galaxies at the cost of a density highly correlated with

redshift. The next sample consists of Luminous Red Galaxies (LRG), which are large elliptical galaxies that are strong tracers of the underlying matter (Zhou et al. 2023). We split the LRGs into three redshift bins between $0.4 < z < 1.1$. The final sample contains Emission Line Galaxies (ELG), which are largely blue, spiral galaxies (Raichoor et al. 2023). We split these into two redshift bins between $0.8 < z < 1.6$. The LRG and ELG samples overlap in $0.8 < z < 1.1$, allowing for cross-correlations. All subsamples are labeled as their galaxy type and a number which increases with redshift, e.g. the lowest-redshift BGS sample is BGS1. The properties of these subsamples are displayed in Table 1.

3. DEFINING MULTIPLETS

3.1. Identifying Multiplets

Multiplets are sets of galaxies separated from each other by a maximum linking length ℓ and identified similarly to the friends-of-friends (FOF) algorithm (More et al. 2011). We begin by finding all pairs of galaxies separated by ℓ or less. We then identify multiplets as the connected components of the resulting pairs using a Union Find algorithm, setting no limit on the physical size or number of members of a group. To account for the differing scales of line-of-sight (LOS) and transverse uncertainty, largely from redshift-space distortions, we split the linking length into a transverse component ℓ_p and a LOS component $\ell_{\parallel}$.

Our goal is to measure tidal shear with the highest signal-to-noise ratio (SNR), rather than to identify a consistent set of physical objects. The scales at which multiplets are defined therefore vary between samples, with denser samples and higher clustering favoring smaller linking scales. We do not expect variations within the $\ell \sim 1 h^{-1} \text{ Mpc}$ regime to affect the scale dependence of the signal, and expect only a small impact on its overall amplitude.

To identify the optimal linking lengths, we measured multiplet-galaxy alignment across a range of scales in each sample (described later in Section 5). The signal was averaged in six logarithmic bins between projected separations of 10 and $50 h^{-1} \text{ Mpc}$ and the errors are the standard error from measuring in different sky regions; we found no improvement from a full jackknife.

We sampled linking scales of both ℓ_p and $\ell_{\parallel}$ in steps of $0.5 h^{-1} \text{ Mpc}$, up to $\ell = 8 h^{-1} \text{ Mpc}$ and selected the smallest ℓ which produced a measurement within 1% of the peak SNR, on the basis that adding a large number of additional multiplets is not warranted for a marginal gain in signal. We perform this test on a subset of the area in all samples, using the minimum area necessary to obtain smooth results with respect to ℓ . This is more resource efficient and ensures we are not sensitive to a SNR spike in the overall measurements.

This SNR scales as A/σ , where σ is the per-multiplet scatter, the same quantity that governs the final measurement significance, making it a well-motivated cut. The scale that maximizes this SNR need not coincide with the one that maximizes the signal amplitude A , since a larger linking scale alters the multiplet population and can also increase σ ; in practice, however, the two peaks fall at the same scale for the samples considered here. Additionally, since the signal appears to be constant over a range of ℓ , our chosen scales are the ones which minimize noise.

For LRG, we restricted the measurement to a patch defined by $130 < \text{RA} < 150$ and $0 < \text{Dec} < 20$, in degrees. For

the sparser ELG sample, we used the full catalog and increase the measured scale to $100 h^{-1}\text{Mpc}$. We found $\ell_p = \ell_{\parallel} = 5 h^{-1}\text{Mpc}$ to be optimal for all redshift bins in the LRG sample, and $\ell_p = 7 h^{-1}\text{Mpc}$, $\ell_{\parallel} = 3 h^{-1}\text{Mpc}$ for ELGs. This agreement across redshift bins within the samples is expected, given their near-constant comoving number density.

Conversely, the density of the magnitude-limited BGS sample drops substantially toward higher redshifts. We tested a range of linking scales in a subset of the footprint for each redshift bin, using 800, 800, 1500, 1500, and 3000 deg^2 for BGS1 through BGS5 respectively, with the largest areas reserved for the sparser high-redshift bins and the smallest for the two densest, lowest-redshift bins. Across all BGS redshift bins, we found no significant benefit to adopting different scales for ℓ_p and $\ell_{\parallel}$ and set the two equal. The results are shown in Figure 1. Especially for lower redshifts, the scale choice does not have a dramatic impact on the signal amplitude above $1.5 h^{-1}\text{Mpc}$.

The existence of a peak in SNR across varying linking scales reflects a competition between two effects. Increasing the linking scale preferentially builds larger objects, which trace the tidal field more strongly (a known effect for galaxies and halos), and galaxies with larger separations are less likely to influence each other's positions, which washes out external tidal information. Beyond a certain scale, however, the linking becomes loose enough to associate galaxies with progressively weaker physical connection, diluting the coherent orientation signal. The balance of these effects for a given galaxy sample sets the optimal scale. Notably, the linking lengths we find are considerably larger than those adopted in Lamman et al. (2024b) of $\ell_p = 1 h^{-1}\text{Mpc}$ and $\ell_{\parallel} = 6 h^{-1}\text{Mpc}$, suggesting that large-scale structure imprints a coherent orientation across more loosely separated galaxies than previously utilized.

It is also somewhat surprising that we recover $\ell_p = \ell_{\parallel}$ for every sample except ELG. We attribute this to the fact that the optimal scales are large enough that redshift-space distortions are negligible, removing the primary motivation for an elongated LOS linking length. The exception is the ELG sample, where we find a transverse scale more than twice the LOS scale (7 versus $3 h^{-1}\text{Mpc}$). We interpret this difference as a consequence of sample variance; ELGs are the sparsest and most diffuse of our samples. We nonetheless retain the recovered values, as our aim here is to obtain the highest-SNR measurement possible rather than a physically uniform multiplet definition across samples.

Ideally, future analyses will similarly identify the optimal scales for each sample. We note that, with these samples, the optimal linking scale roughly scales with galaxy number density as $\ell \propto \ell_0 - \frac{1}{4} \log n$.

This selection is robust to cosmology: because we quote all separations in $h^{-1}\text{Mpc}$, ℓ does not depend on H_0 , and only negligibly on the shape of the distance-redshift relation.

3.2. Multiplet Shapes

The orientation of each multiplet is determined from the projected, plane-of-sky positions of its member galaxies. Using projected positions mitigates the effect of small-scale RSD on the inferred orientation. Following Lamman et al. (2024b), the multiplet orientation Θ is the angle of the N -member-averaged




FIG. 1.— The results of measuring multiplet alignment in a subset of the BGS catalog (defined in Table 1) as a function of linking scale, the maximum separation between two galaxies for them to be considered part of the same multiplet. The choice of linking scale has a minimal effect on the signal amplitude (top panel) above $\ell \sim 1.5 h^{-1}\text{Mpc}$. We select the scale based on the total SNR (bottom panel), which captures where most of the alignment information lies. This is especially important for BGS, whose number density varies strongly with redshift. The chosen linking scale for each redshift bin is marked by a vertical dashed line and, as expected, increases as the number density decreases.

complex position relative to the multiplet centroid,

$$\epsilon_{\text{multiplet}} = \frac{1}{N} \sum_{i=1}^N r_i e^{2i\theta_i} = a + bi, \quad \Theta = \frac{1}{2} \arctan \frac{b}{a}, \quad (1)$$

where r_i and θ_i are the projected distance and position angle of member i relative to the centroid. This is a “stick” model: each multiplet is characterized by orientation alone, with no axis-ratio information.

The error of a shape alignment measurement scales with the ‘shape noise’ of the shape sample, conventionally taken as the variance of one component of the complex ellipticity. For multiplets modeled as sticks, the orientation is a unit spin-2 phasor $e^{2i\theta}$ with θ uniformly distributed under the null hypothesis, so $\langle \cos^2 2\theta \rangle = \langle \sin^2 2\theta \rangle = 1/2$ and the per-component shape noise is $\sigma_s = 1/\sqrt{2}$, about three times higher than many galaxy samples.

To potentially bring this down, we try using the full shape information of multiplets. The full ellipticity (E_1, E_2) projected perpendicular to the $\hat{z}$ line of sight can be computed from the second moments Q of member positions about the multiplet center:

$$E_1 + iE_2 = \frac{Q_{xx} - Q_{yy} + 2iQ_{xy}}{Q_{xx} + Q_{yy}}. \quad (2)$$

This is the standard second-moment definition used to compute ellipticities from particle distributions (Bartelmann & Schneider 2001), and reduces to the orientation-only ‘stick’ ($|E| = 1$) when a multiplet has only two members. For a multiplet of N members drawn from the same orientation distribution, the




FIG. 2.— Comoving number densities of galaxies (solid) and the multiplets identified within them (dashed) as a function of redshift, for the BGS, LRG, and ELG samples. The horizontal bars mark the redshift extent of each sample bin. The multiplet density tracks the galaxy density approximately linearly across all samples, remaining near $1/4$ of the galaxy density for BGS and LRG, $1/6$ for ELG.

ellipticity is the mean of N such phasors, so its per-component variance scales as $1/N$ and the shape noise as $\sigma_s \propto 1/\sqrt{N}$.

However, this also decreases the measured amplitude of multiplet alignment. The total SNR of a shape-tracer measurement scales with alignment amplitude A and shape noise σ_s as A/σ_s . When measuring alignment with the full ellipticity, we find the amplitude decreases faster than the shape noise: averaging member orientations within a multiplet washes out the alignment signal more than it suppresses the noise. The SNR is not improved by using the full ellipticity, so we retain the stick model.

3.3. Multiplet Demographics

Here we examine the number densities of multiplets, along with their sizes, shapes, and number of members. Although the physical nature of a multiplet is not directly relevant to the large-scale alignment measurement, their characterization is helpful in building physical intuition and astrophysical interpretations.

To estimate number densities, the comoving volume of each sample was computed by treating its galaxies as tracers of the survey footprint. The RA/Dec positions were pixelized on a HEALPix grid ($n_{\text{side}} = 64$), and the number of occupied pixels multiplied by the pixel area gives the effective solid angle Ω . Combined with the comoving distances at the bin's redshift edges, this yields the shell volume. The number density $n(z)$ for both galaxies and multiplets was then computed on a finer redshift grid within each bin using this same footprint estimate, ensuring the summary volumes and densities share consistent convention.

The number densities of multiplets and the galaxy samples in which they were identified are shown in Table 1 and Figure 2. In BGS, the multiplet number density is consistently close to $1/4$ of the galaxy density, despite the linking length varying between the redshift subsamples. For LRG and ELG, the multiplet fraction is somewhat lower, roughly $1/5$ to $1/4$. Across all samples, the multiplet number density scales only

slightly more steeply than linearly with galaxy number density. This is notable given that the number of candidate neighbors available to form multiplets scales as n^2 , which would naively suggest a much steeper relation. The near-linear scaling instead reflects the clustering of galaxies: the close pairs that seed multiplets reside in overdense regions, which are already populated at low mean density, so increasing the overall density predominantly adds galaxies in underdense regions that do not form new links. This effect is compounded by our Union Find definition, which places no cap on membership. In the densest regions, added galaxies tend to grow existing multiplets rather than nucleate new ones, further flattening the relation.

We also visualize typical multiplet shapes in Figure 3. For each sample and redshift bin, we constructed a stacked image of multiplet shapes as follows. Each multiplet's member galaxies were first rotated into a frame aligned with that multiplet's own orientation axis, defined in Section 3.2. The rotated members of every multiplet in a given subsample were then pooled into a single point cloud of comoving separations from the multiplet centroid. The two-dimensional density of this point cloud was estimated via a Gaussian-smoothed histogram in Cartesian coordinates and resampled onto a polar grid for display. Density was mapped to opacity at fixed hue, shown out to a fixed comoving radius of $4 h^{-1} \text{Mpc}$. The average number of galaxies per multiplet for each sample is shown below its corresponding panel. Histograms of the number of multiplet members for the DESI DR1 sample are shown in Lamman et al. (2024b), and are qualitatively unchanged besides a larger linking length resulting in a larger average number of members.

The majority of multiplets are pairs, which dominate the stacked shapes and give them their characteristic elongated, stick-like appearance. However, galaxies in multiplets with more than two members also tend to align closely with the multiplet's primary axis. The gap in the center of each panel is set by the minimum physical separation of pairs resolved in the survey, which grows with redshift. Many multiplets extend well beyond the scale of a typical halo or cluster ($R_{200} \sim 1\text{--}2 h^{-1} \text{Mpc}$ for massive clusters). This, combined with their distinct elongation, suggests that the typical multiplet more closely resembles a filamentary structure, though groups and clusters are also captured in the multiplet catalogs. We also emphasize that the physical meaning of a multiplet varies somewhat across galaxy samples and redshift bins, depending on the number density of the parent sample and chosen linking length.

4. MULTIPLET ALIGNMENT

4.1. Alignment Estimators

We measure the alignment of multiplet orientations with respect to a density tracer ($\mathcal{E}_+$), and the autocorrelation of multiplet orientations with each other ($\mathcal{E}_{++}$). The former follows Lamman et al. (2024b); we summarize the relevant points here.

The cross-correlation estimator averages the projected ellipticity $\epsilon_+(j|i)$ of each shape $i \in S$ relative to each tracer $j \in D$,

$$\mathcal{E}_+(r) = \frac{1}{N_{\text{pairs}}} \sum_{i \in S, j \in D} \epsilon_+(j|i) = \frac{S_+ D(r)}{D D(r)}, \quad (3)$$

where ϵ_+ quantifies the multiplet orientation relative to the




FIG. 3.— Stacked visualization of the average multiplet shape in each galaxy sample and redshift bin, projected onto the plane of the sky. Member galaxies are rotated into each multiplet’s own orientation frame before stacking, so the common elongation is preserved; density is mapped to opacity. All panels are shown out to a fixed comoving radius of $4 h^{-1} \text{Mpc}$ from the multiplet center, and the average number of members per multiplet $\bar{N}$ is shown below each panel. Because most multiplets are pairs, the stacks are dominated by elongated, stick-like shapes. However, even galaxies in $N > 2$ multiplets fall closely along the multiplet’s primary axis. The central gap reflects the minimum resolved pair separation, which increases with redshift. The overall size grows toward sparser, higher-redshift samples with larger linking lengths.

axis connecting the multiplet and the tracer, θ , as

$$\epsilon_+ = \cos 2\theta_i. \quad (4)$$

A positive value indicates tidal alignment. S_+D represents the number of alignment-weighted shape-tracer pairs and DD the number of tracer-tracer pairs. This is related to the shape-density correlation function ξ_{g+} , a generalized form of the Landy-Szalay estimator (Mandelbaum et al. 2006),

$$\xi_{g+} = \frac{S_+D - S_+R_D}{R_S R_D} \approx \frac{S_+D}{R_S R_D}. \quad (5)$$

The key difference is that ξ_{g+} is not averaged on a per-pair basis. The two estimators are related by

$$\xi_{g+}(s) = \frac{DD}{R_S R_D} \mathcal{E}_+(s) = (1 + \xi(s)) \mathcal{E}_+(s). \quad (6)$$

The signal is projected along the line of sight up to a maximum separation Π_{max} . Following Lamman et al. (2025), we adopt a variable Π_{max} that depends on the bin of transverse separation: $\Pi_{\text{max}}(R_{\text{bin}}) = [8 + \frac{2}{3} R_{\text{bin}}] h^{-1} \text{Mpc}$. This optimizes signal-to-noise and captures comparable information to 3D estimators, while remaining simpler to measure and model. We do not apply the full Gaussian LOS weighting introduced in that work, as the marginal gain in signal-to-noise does not justify the additional computational cost for this analysis.

In this work we additionally measure the multiplet orientation autocorrelation, $\mathcal{E}_{++}$. For a pair of multiplets i and j , we project each orientation onto their connecting separation vector and correlate the two projected components as a function

of transverse separation R ,

$$\mathcal{E}_{++}(R) = \frac{1}{N_{\text{pairs}}} \sum_{i,j \in S} \epsilon_+(j|i) \epsilon_+(i|j) = \langle \cos 2\theta_i^{(R)} \cos 2\theta_j^{(R)} \rangle, \quad (7)$$

where $\theta_i^{(R)}$ and $\theta_j^{(R)}$ are the orientation angles of the two multiplets relative to the vector connecting them. This is the orientation-only analog of the projected shape-shape correlation w_{++} commonly measured for galaxy ellipticities (Singh & Mandelbaum 2016; Troxel & Ishak 2015), with the multiplet ellipticity modulus folded into the amplitude parameter introduced in Section 4.2. As with $\mathcal{E}_+$, the signal is projected along the LOS using the same variable Π_{max} .

Unlike the cross-correlation between multiplet orientations and the density field, the orientation autocorrelation does not require a random catalog. In the cross-correlation, randoms serve to normalize the summed shears by the available pair counts and to subtract the spurious alignment induced by the survey footprint. Neither role applies here. Both quantities entering the autocorrelation are shape orientations, which have zero mean and carry no monopole sourced by the mask.

4.2. Modeling

We model the multiplet alignment signal by assuming a linear relationship between orientations and the large-scale tidal field, following the nonlinear alignment (NLA) framework (Bridle & King 2007) as applied to multiplets in Lamman et al. (2024b). The projected ellipticity is related to the traceless tidal tensor $T_{\alpha\beta}$ by a single amplitude parameter τ ,

$$\epsilon_{\alpha\beta} = \tau \left(T_{\alpha\beta} + \frac{1}{2} T_{zz} \right), \quad (8)$$

Sample	z range	Volume [Gpc ³ h ⁻³]	ℓ_p [h ⁻¹ Mpc]	$\ell_{\parallel}$ [h ⁻¹ Mpc]	N_{galaxies} [10 ⁶]	$N_{\text{multiplets}}$ [10 ⁶]	$-\tau$ [10 ⁻²]	A_{IA}
BGS1	0.01 - 0.1	0.04	1.5	1.5	1.19	0.33	4.05 ± 0.45	8.64 ± 0.96
BGS2	0.1 - 0.2	0.25	2	2	2.74	0.75	6.07 ± 0.39	10.67 ± 0.69
BGS3	0.2 - 0.3	0.61	2.5	2.5	2.49	0.65	7.60 ± 0.17	11.09 ± 0.25
BGS4	0.3 - 0.4	1.07	3.5	3.5	1.24	0.30	9.65 ± 0.49	11.92 ± 0.61
BGS5	0.4 - 0.5	1.58	5	5	0.36	0.08	11.42 ± 1.76	12.19 ± 1.88
LRG1	0.4 - 0.6	3.44	5	5	1.07	0.24	11.80 ± 0.09	11.37 ± 0.09
LRG2	0.6 - 0.8	5.41	5	5	1.64	0.37	13.62 ± 0.27	10.41 ± 0.21
LRG3	0.8 - 1.1	11.3	5	5	1.83	0.35	14.66 ± 0.38	9.04 ± 0.23
ELG1	0.8 - 1.1	11.3	7	3	2.85	0.50	5.51 ± 0.24	3.34 ± 0.15
ELG2	1.1 - 1.6	24.8	7	3	3.84	0.56	7.86 ± 0.17	3.68 ± 0.08

TABLE 1. Properties of our multiplet samples and the galaxy catalogs they were identified in. The last two columns are our reported values for the alignment strength of the multiplets, as measured by the multiplet-galaxy cross-correlation. Shown is both our parametrization of the tidal response, τ , and the commonly-used alignment amplitude for the intrinsic alignments of galaxies, A_{IA} . Errors on both of these scale-independent amplitudes come from the variance in estimating them in 12 logarithmically spaced R bins between 20-100 $h^{-1}\text{Mpc}$.

where τ encodes the response of multiplet orientations to the tidal field, including any contribution from the full-shape information and any systematic misalignment, assuming neither introduces scale dependence at large separations. For the cross-correlation, this yields the model prediction derived in (Lamman et al. 2024b)

$$\mathcal{E}_+(R) = \frac{-\tau}{2\Pi_{\text{max}} + \bar{w}_p} \int K dK \mathcal{J}_2(K, R) \mathcal{P}_{\Pi}(K), \quad (9)$$

where $\mathcal{J}_2$ is the bin-averaged second-order Bessel function of the first kind, $\mathcal{P}_{\Pi}(K)$ is the LOS-projected matter power spectrum weighted by the tidal factor K^2/k^2 , and $\bar{w}_p$ is the bin-averaged projected cross-correlation between the multiplet and tracer catalogs. The power spectrum enters as $b_g P_{mm}(k)$, with b_g the galaxy bias of the tracer sample.

We characterize the alignment amplitude with τ rather than the more common A_{IA} used in most weak-lensing-related works (Bridle & King 2007; Joachimi et al. 2011; Johnston et al. 2019), as τ is a more natural characterization of the pure tidal response, although we report both here. The two are related by

$$A_{\text{IA}}(z) = -\frac{\tau}{C_1} \frac{D(z)}{\rho_m(z)}, \quad (10)$$

where $C_1 = 5 \times 10^{-14} M_{\odot}^{-1} h^{-2} \text{Mpc}^3$ (Brown et al. 2002), ρ_m is the matter density, and $D(z)$ is the growth factor normalized such that $\bar{D}(z) = (1+z)D(z)$ is unity at $z=0$. Note that here we have changed the $D(z)$ normalization from Lamman et al. (2024b), which normalized at matter domination, to match the majority of recent works. Since A_{IA} contains the growth factor while τ does not, the assumed cosmology enters directly into A_{IA} and is another reason to adopt τ as our primary parameterization.

The autocorrelation $\mathcal{E}_{++}$ follows from the same framework with two modifications. First, because both fields being correlated are shapes rather than one shape and one density tracer, the amplitude enters quadratically: $\mathcal{E}_{++} \propto \tau^2$, and the galaxy bias does not appear, since neither field is a density tracer. The relevant power spectrum is therefore $P_{mm}(k)$, with the tidal factor entering squared, $(K^2/k^2)^2$. Second, the projection geometry changes. Correlating two spin-2 (shape) fields, rather than a shape and a scalar, replaces the $\mathcal{J}_2$ kernel of the cross-correlation with the sum $\mathcal{J}_0 + \mathcal{J}_4$. This is the standard projection for shape–shape correlations (Singh & Mandelbaum 2016; Hervas Peters et al. 2025). The bin-averaged Bessel kernels are defined as in Lamman et al. (2024b) for $\mathcal{J}_2$, with

$\mathcal{J}_0$ and $\mathcal{J}_4$ following the same convention. The $\mathcal{J}_0$ term makes $\mathcal{E}_{++}$ more isotropic in the (r_p, Π) plane than $\mathcal{E}_+$, while the $\mathcal{J}_4$ term retains the transverse dependence. We confirm that these bin-averaged transforms converge with respect to the upper integration limit. The model prediction is then

$$\mathcal{E}_{++}(R) = \frac{\tau^2}{2\Pi_{\text{max}} + \bar{w}_p} \int K dK 2[\mathcal{J}_0(K, R) + \mathcal{J}_4(K, R)] \mathcal{P}_{\Pi}^{(2)}(K), \quad (11)$$

where $\mathcal{P}_{\Pi}^{(2)}$ carries the squared tidal factor $(K^2/k^2)^2$ and $\bar{w}_p$ is the bin-averaged multiplet–multiplet projected correlation function. We do not include a factor of $\frac{1}{2}$ in this expression (for a per-multiplet quantity) in order to be consistent with the $\mathcal{E}_+$ estimator, which is defined as an average over pairs.

Crucially, if τ is determined from the better-constrained cross-correlation, Equation 11 is a parameter-free prediction of the autocorrelation amplitude: there is no free normalization, and the galaxy bias does not enter. If, instead, τ is measured from the autocorrelation, the cross-correlation provides an independent measurement of the galaxy bias. In the following sections, we adopt the τ value from the $\mathcal{E}_+$ fit and overlay the resulting τ^2 prediction on the measured $\mathcal{E}_{++}$ as a consistency check of the alignment framework. We additionally fit τ directly to $\mathcal{E}_{++}$, compare the two amplitude values, and check consistency with the predicted galaxy bias.

5. MEASUREMENTS

5.1. Measurement Method

We measure both alignment estimators of Section 4.1 with the pipeline introduced in Lamman et al. (2024b), updated for the larger DR2 samples. The code is publicly available in the `spec-IA`¹ repository. We summarize the procedure here and describe the changes made for this work.

We compute each projected estimator in bins of transverse separation, where larger bins have larger Π_{max} , as defined in Section 4.1. The measurement is carried out independently in N_{reg} sky regions of equal multiplet count, divided in right ascension and declination. Within each region, the multiplet orientations are measured relative to the full tracer sample, or full multiplet sample for autocorrelations. The average signal and errors are then estimated from these regions. We increase

¹ github.com/cmlamman/spec-IA

the number of regions relative to Lamman et al. (2024b), using $N_{\text{reg}} = 2500$ for BGS2 (which contains the largest number of multiplets), $N_{\text{reg}} = 2000$ for the other BGS samples, and $N_{\text{reg}} = 1000$ for LRG and ELG. The larger count for BGS is driven by its much higher density, both to keep the per-region pair count tractable in memory and because the denser sample supports a finer regional division at fixed statistical precision. Alongside this, we made minor optimizations to the pair-counting code; the measured signal is unchanged from the original implementation.

Errors are estimated by bootstrap resampling of the regions: regions are drawn with replacement $N_{\text{boot}} = 10^5$ times and the per-bin uncertainty is the standard deviation of the resampled mean. The same resampling defines the covariance,

$$\mathbf{C} = \text{Cov}[\bar{\mathbf{d}}^{(b)}]_{b=1}^{N_{\text{boot}}}, \quad (12)$$

where $\bar{\mathbf{d}}^{(b)}$ is the mean over the b -th resampling, so error bars and covariance follow from a single procedure.

For some analysis, we combine the measurements for each tracer across their redshift subsamples. For this, we use an amplitude-based weighting: each subsample is divided by its best-fit alignment amplitude τ before the inverse-variance combination. This weighting has little effect for the LRG and ELG samples, whose subsamples have similar amplitudes, but modestly improves the combined BGS signal-to-noise. When combining, the bins are disjoint and treated as independent. They are combined at the covariance level: the combined inverse covariance is $\mathbf{C}_{\text{comb}}^{-1} = \sum_z \mathbf{C}_z^{-1}$ and the combined signal is $\bar{\mathbf{d}}_{\text{comb}} = \mathbf{C}_{\text{comb}} \sum_z \mathbf{C}_z^{-1} \bar{\mathbf{d}}_z$. Inverse covariances are debiased with the factor $(N_{\text{reg}} - N_{\text{bin}} - 2)/(N_{\text{reg}} - 1)$ before combination (Hartlap et al. 2007).

We validated the pipeline against an independent implementation built on TREECORR (Jarvis et al. 2004), finding agreement in the measured signal. We nonetheless retain our own code for the production measurements, as TREECORR does not natively support the variable $\Pi_{\text{max}}(R_{\text{bin}})$ binning or the three-dimensional pair selection used here, and offers no sufficient speed advantage once these are imposed.

5.2. Cross-correlations

For the cross-correlation $\mathcal{E}_+$, we compute the projected orientation of each multiplet relative to each tracer as described in Section 4.1. Unless otherwise stated, the tracer catalog is the same catalog used to identify the multiplets; the exception is the overlapping LRG and ELG region at $0.8 < z < 1.1$, where we cross-correlate the two samples.

For every measurement, we also measure the multiplet orientations relative to DESI’s random catalogs, which are constructed to match the survey geometry, and subtract this from the signal. The randoms subtraction differs from Lamman et al. (2024b). Previously, the mean random signal was subtracted from the mean measured signal after averaging over regions. Here we instead subtract the randoms on a per-region basis before averaging. The two procedures yield the same combined signal, but the per-region subtraction removes the coherent, geometry-sourced scatter between regions before the variance is estimated. The resulting errors are more accurate and smaller at large separations, where the geometric term contributes most to the region-to-region variance.

The resulting cross-correlation measurements and NLA model fit are shown in Figure 4. Across all samples, we see a significant improvement from the DR1 measurements

presented in Lamman et al. (2024b). This is partially due to improved measurement methods, but mostly due to more data and adopting a larger linking length when identifying multiplets. This especially improved the ELG signal, which had the largest increase in galaxy density from DR1 and, as the most diffuse sample, had the most to gain from a larger maximum linking length. Compared to DR1, the number of BGS and LRG galaxies roughly doubled, while ELG increased by a factor of about 2.5. The number of multiplets in each sample, largely due to the new definition, increased by factors of 3 for BGS, 9 for LRG, and 25 for ELG.

Measurements from BGS and ELG are consistent with the NLA model in the large-scale linear regime, generally agreeing over the scales it is expected to be valid, $R > 10h^{-1}\text{Mpc}$ (Chisari 2025). However, Lamman et al. (2024b) found NLA to break down for LRGs starting below $30h^{-1}\text{Mpc}$, which is similarly reflected in DR2 LRGs. This could be because density is correlated with alignment strength in ways that NLA does not capture (Ichikawa et al. 2026). In particular, the observed shape field carries a density weighting factor of $(1 + \delta)$, which our $1/(1 + \xi)$ normalization (Equation 6) can only partially account for. This increases the minimum valid NLA scales for samples with higher clustering (Shi et al., in prep), and would be especially pronounced for highly biased tracers like multiplets.

Within the BGS sample, signal amplitude is highly correlated with redshift. This is largely due to the clustering differences between redshift bins, as the galaxy bias increases significantly when fainter galaxies drop out of the more distant bins. Because of this, the maximum linking length set for identifying multiplets also changes across the bins (Table 1). This difference in clustering and multiplet definition is accounted for in the alignment amplitudes discussed in Section 7.3; this plot should not be used to infer any redshift dependence. Conversely, the LRG and ELG catalogs have a more constant number density (Figure 2) and identical multiplet definitions within each sample, so the amplitude of the correlations are similar across redshift bins.

The overlap of LRGs and ELGs in the redshift range $0.8 < z < 1.1$ provides an interesting opportunity to directly compare the alignment of multiplets made with galaxies of different morphologies, so the cross-correlation of ELG multiplets with LRG tracers is also shown in Figure 4. This is discussed further in Section 7.3.

Figure 6 shows the average multiplet alignment made with BGS, LRG, and ELG. They are a combination of the measurements in all redshift bins for each tracer, with errors taking into account the difference in amplitude across bins, as described above. Differences in amplitudes on this plot are from both clustering variations between the samples and differences in the alignment strength of multiplets. These represent our cleanest measurements of large-scale tidal shear with each tracer. The correlation matrix of the combined LRG signal is characteristic of all measurements and is shown in Figure 7. The diagonal of the underlying covariance reproduces the standard errors in Figure 6. Adjacent R bins are positively correlated, especially at large separations. More details of this covariance estimation, the impact of region size, and detection significance are discussed in Section 7.1.

5.3. Autocorrelations

The multiplet orientation autocorrelation $\mathcal{E}_{++}$ (Section 4.1) is measured with the same infrastructure as $\mathcal{E}_+$, differing only in the quantity correlated and in the absence of a random cat-




FIG. 4.— Correlations between the projected shape of multiplets and the surrounding matter field, as traced by the same galaxy samples they were identified in. This is measured as function of projected separation between multiplets and tracers, R , and with an R -dependent LOS distance, described in Section 4.1. Within a given R bin, a larger value corresponds to stronger tidal alignment. The signals are compared to the data-normalized prediction of the nonlinear alignment model (NLA), shown in dotted lines. The BGS measurements (left) have a high correlation with galaxy clustering, which increases with redshift across the BGS sample. Measurements from the LRG sample (center) are more consistent across redshift bins as the sample is designed to have a constant co-moving number density. Measurements made with ELGs (right) show the most striking improvement over DR1, and are especially notable given that the measured alignment of individual ELGs is consistent with 0. The right panel also displays the cross-correlation of ELG multiplets with LRG tracers. The increased signal demonstrates the dependence of signal amplitude on galaxy clustering, however it is still half of the amplitude of LRG multiplets measured relative to the same tracer sample. This demonstrates the tendency of ELG multiplets to fall in less dense regions and is discussed more in Section 7.3.




FIG. 5.— The autocorrelation of multiplet shapes in the BGS, LRG, and ELG samples, measured as a function of projected separation R and the same variable LOS distance used in Figure 4. The multiplet autocorrelation is a weaker signal than the cross-correlation because shape-shape correlations are intrinsically weaker, and there are far fewer multiplets than individual galaxies, which are used in the cross-correlation. This is especially apparent in less dense regions, such as ELGs, where the signal is marginally detectable. The highest redshift BGS bin is similarly noisy and has been removed from this figure for clarity. The measurements of LRG and ELG are averaged over the redshift bins within each sample, and all measurements are compared to the NLA model prediction made with two amplitudes. The first is directly fit to the measured autocorrelation, $\mathcal{E}_{++}$, shown in a faint solid line. The second is the amplitude fit to the cross-correlation, $\mathcal{E}_{+}$, shown in a dashed line. The model fit to the autocorrelation amplitude is systematically smaller, likely due to an underestimated galaxy bias (Section 5.3).

alog. Rather than correlating a multiplet orientation with a tracer position, we take pairs of multiplets, project each multiplet's orientation onto the separation vector connecting the pair, and average the product of the two projected components. The line-of-sight pair selection $\Pi_{\max}(R_{\text{bin}})$, the division into sky regions of equal multiplet count, and the estimation of errors and covariance from the region-to-region scatter are all unchanged from the cross-correlation.

The resulting autocorrelation measurements are shown in Figure 5. Shape-shape correlations are generally weaker than shape-density cross-correlations, and this effect is more pronounced with multiplets than individual intrinsic alignments since there are fewer shape tracers (multiplets) than density tracers (galaxies). The ELG measurements especially suffer

from this, where the density of multiplets are low, and we find only a marginal detection of the ELG autocorrelation. The signal for all samples is shown in Figure 5, where the signal for LRGs and ELGs has been averaged over the redshift bins.

They are compared to the model predictions normalized with two amplitudes: one directly fit to the measured autocorrelation, $\mathcal{E}_{++}$, and one fit to the cross-correlation, $\mathcal{E}_{+}$. Across all samples, the amplitude fit to $\mathcal{E}_{+}$ is larger and shows evidence of a true tension in LRG and BGS. For LRGs, the amplitudes are consistent for LRG1, but the combined disagreement across the three redshift bins reaches $\sim 5\sigma$. This could be from the assumption of one galaxy bias value across the redshift range, as opposed to an issue with the NLA model. For BGS, we use the independent, magnitude-calibrated bias




FIG. 6.— Intrinsic tidal shear measured with galaxy multiplets in DESI’s DR2 BGS, LRG, and ELG samples. This is the largest-scale detection of its kind. The measurements are made as a function of projected separation R , with a $\Pi_{\max}$ value that increases with R (shown on the top axis). They are the signals in Figure 4, combined across redshift bins, and compared to NLA model predictions. The significance of each signal being detected above 100, 200, and 300 $h^{-1}\text{Mpc}$ is shown on the plot. These significances are computed with a matched filter based on the model shape, downweighting measurements where the correlation is known to be weaker. The high density of the BGS sample and the strong bias of the LRG sample result in high-significance detections, even when the correlation strength drops at large separations. We note that the BGS significance is an upper limit as our covariance may underestimate sample variance (Section 7.1). Although a weaker detection, ELG multiplet alignment is especially notable given that the equivalent signal for individual ELGs is consistent with 0; this represents our best approach for detecting intrinsic shear above redshift $z = 1$.

values from the UCHUU-BGS mocks (see below), but still see disagreement at the 4–8 σ level depending on the adopted magnitude percentile. There is also a large fractional offset in ELG, although this is consistent with its noisy detection. We find no evidence that this is a result of the model itself. The Bessel-kernel projections and tidal-factor treatment used in $\mathcal{E}_+$ and $\mathcal{E}_{++}$ are unchanged between samples, so the redshift-dependence of the tension is more consistent with an underestimated galaxy bias. Where the discrepancy is worst, LRG3, it would be reconciled with an increase in the galaxy bias of 25%.

The high density of multiplets in BGS allows for significantly cleaner measurements of the multiplet autocorrelation. In part due to this, and in part due to the difficulty of estimating galaxy bias for a sample that varies so much across redshift bins, we decided to use these correlations to independently estimate the galaxy bias directly from the samples, as described below.

5.4. Galaxy Bias

The inferred alignment amplitude τ is inversely proportional to the linear galaxy bias, b_g , so accurate bias estimates are essential, particularly when comparing measurements across redshift bins with different sample properties. For the LRG and ELG samples we adopt the linear bias values from DESI’s fiducial cosmological analyses of $b_{\text{LRG}} = 1.99$ and $b_{\text{ELG1}} = 1.18$, $b_{\text{ELG2}} = 1.40$ (Mena-Fernández et al. 2025; Vaisakh et al. 2026). For BGS, the situation is more complicated, and we describe our approach in this section. We compare to reference mocks but ultimately take advantage of the fact that $\mathcal{E}_+$ depends on b but $\mathcal{E}_{++}$ doesn’t, and estimate it from our samples directly.

The BGS BRIGHT sample used in this work spans $0.01 < z < 0.5$, extending well beyond the redshift range $0.1 < z < 0.4$ used in DESI’s standard cosmological BGS analyses. Unlike the LRG and ELG samples, which are designed to maintain

a roughly constant comoving number density with redshift, BGS is flux-limited: the $r < 19.5$ magnitude cut preferentially removes fainter, less-biased galaxies at higher redshifts. As a result, both the effective luminosity and the linear bias of the BGS sample evolve strongly across our five redshift bins.

We first estimate b_g using the fit to the UCHUU-BGS mock catalogs from Fernández-García et al. (2026), who measure the linear bias of BGS BRIGHT as a function of absolute magnitude threshold. They parameterize the bias of a volume-limited sample of galaxies with M_r brighter than a given threshold as:

$$b(< M_r) = B_0 + B_1 \times 10^{B_2(M_r^* - M_r)/2.5}, \quad (13)$$

with best-fit parameters $B_0 = 1.087$, $B_1 = 0.196$, $B_2 = 1.12$, and $M_r^* = -20.44$ for the DR2 BGS BRIGHT sample. This fit is calibrated over $-22 \leq M_r \leq -20$, but the functional form is monotonic and well-behaved outside this range.

Because our BGS bins are flux-limited rather than volume-limited, we cannot directly apply Equation 13 at a single threshold. Instead, we approximate the effective bias in each bin by identifying M_r^{eff} as a high percentile of the observed M_r distribution in that bin. This approximation becomes less accurate at low redshift, where the flux limit permits a wider range of intrinsic luminosities. We compute the bias using 99.9th, 99.99th, and 100th (brightest galaxy) percentiles; results are shown in Figure 8. The estimates are relatively stable at low z and diverge modestly at high z , reflecting the extrapolation into the bright tail.

The multiplet autocorrelation, $\mathcal{E}_{++}$, provides an alternative estimate of b_g that is independent of the mock calibration. As described in Section 4, the cross-correlation amplitude scales as τb_g while the autocorrelation amplitude scales as τ^2 . Taking the ratio for each redshift bin i ,

$$b_g^{(i)} = \frac{|\tau_+^{(i)}|}{\sqrt{|\tau_{++}^{(i)}|}}. \quad (14)$$

These estimates and their uncertainties are also shown in Figure 8.

As an independent cross-check, we predict the relative bias between bins using our own projected two-point correlation function measurements, w_p . In the linear regime, $b_g \propto \sqrt{w_p}/D(z)$, where $D(z)$ is the linear growth factor. We use measurements of the average of our projected correlation function, $w_p^{(i)}$, over R bins of $10 - 50 h^{-1}\text{Mpc}$. In Figure 8, this relation $b(z)$ is shown, scaled to the average multiplet estimate. The shape of $b(z)$ here depends only on our own w_p measurements, and requires no assumption about the alignment amplitude τ or its evolution.

Figure 8 compares all three methods. At low redshift, the multiplet autocorrelation and the UCHUU-BGS fit agree to within $\sim 5\%$. At higher redshift, the UCHUU-BGS estimates rise more steeply than the autocorrelation-derived values, with the discrepancy sensitive to the extrapolation of Equation 13 into the bright end of the luminosity distribution. $b(z)$, normalized to the mean of the autocorrelation estimates, provides a smooth interpolation that lies within the uncertainties of the autocorrelation measurements.

We adopt these per-bin autocorrelation-derived values as our fiducial bias estimates. This choice has three advantages: (1) it is measured from the same data used for the alignment signal, avoiding any mismatch in sample selection or cosmology relative to a mock calibration; (2) each per-bin estimate is independent of τ and its potential redshift evolution, which is a particular concern for BGS given the strong sample-composition changes across our redshift range; and (3) the estimates are consistent with the independent clustering-based shape from w_p and $D(z)$, providing internal validation.

6. MOCK CATALOGS

We reproduce our measurements with mock catalogs built from the ABACUSUMMIT N -body simulations (Maksimova et al. 2021), the same suite used to validate DESI’s clustering analysis. In Lamman et al. (2024b), multiplet alignment measured in these mocks was found to reproduce the scale dependence of the real signal, as expected if the alignment traces the tidal field in the manner described by the NLA framework, while under-predicting its amplitude by about 30% for LRGs. The amplitude offset is a reflection of the simulation’s limited ability to capture the small-scale, higher-order clustering that sources the multiplet orientation. Because the multiplets identified in this work are built with larger linking lengths than in Lamman et al. (2024b) (Section 3), it is not obvious that this behavior should persist.

We measure multiplet alignment in the LRG and ELG mocks; no analogous BGS LSS mocks were available. We averaged over all 25 realizations and compared with the real measurement by taking the ratio of the two averaged over 20 logarithmic bins in projected separation between 10 and $200 h^{-1}\text{Mpc}$. For the three LRG redshift bins the ratio of the measured signal to the mock is 1.20 ± 0.01 , 1.28 ± 0.01 , and 1.33 ± 0.01 for LRG1, LRG2, and LRG3 respectively. These are marginally closer to unity than the ratio of 1.33 reported for LRGs in Lamman et al. (2024b), consistent with the expectation that the larger multiplets used here are less diluted by the nonlinear effects the N -body simulation does not capture. For the two ELG bins, the ratios are 1.97 ± 0.05 and 1.80 ± 0.05 . The larger offset for ELGs follows the same trend: the ELG multiplets are on average smaller and sparser than the LRG multiplets, and so retain a larger contribution




FIG. 7.— The correlation matrix of the combined signal from the three LRG redshift bins, with the identity subtracted to highlight off-diagonal structure. Our quoted significances consider this positive correlation between adjacent bins, which is particularly pronounced at large separations.




FIG. 8.— Estimates of galaxy bias in our BGS redshift samples. The solid lines show predictions from the UCHUU simulations using the faintest P percentiles of galaxies in each bin. The dashed line shows the relative difference in bias between each bin, as predicted from comparing their 2-point clustering, normalized to our estimates from multiplet correlations. These estimates come from comparing the measured amplitude of the cross and auto correlations. They are shown in black, with errors from the variance of determining the amplitudes in independent R bins. The bias values estimated directly from our sample are consistent with the trend expected from 2-point clustering and are the ones we adopt for the rest of this work.

from the small-scale dynamics that the mocks under-represent.

In all cases, the scale dependence of the mock signal remains consistent with the measurement, as in Lamman et al. (2024b), even though the amplitude does not match. We use this agreement in shape to validate the analytic model underlying our BAO analysis in Section 7.4, where it is the scale dependence of the template rather than its overall amplitude that determines the result.

7. RESULTS

7.1. Detection Significance

To compute the significance of our detections above a transverse scale $R_{\min}$, we use the bins with $R > R_{\min}$ and the χ^2 against the null of zero signal,

$$\chi^2 = \bar{\mathbf{d}}^\top \mathbf{C}^{-1} \bar{\mathbf{d}}. \quad (15)$$

Assuming a prior of a positive tidal alignment, this is converted to a one-sided significance. The $\chi^2 \rightarrow \sigma$ conversion assumes Gaussian per-region errors. We validate this with an empirical null distribution of the statistic, built two ways from the same regions: (i) recentering each bootstrap resampling (Eq. 12) to zero signal; (ii) a sign-flip null, randomly flipping the sign of each $\mathbf{d}_i$. Significance is read from the fraction of null realizations exceeding the observed χ^2 . Both nulls agree with the analytic value, confirming Gaussian errors on these scales. With this method, we find a detection significance for the combined LRG signal over $R > 200 h^{-1} \text{Mpc}$ at 3.25σ .

To avoid the dilution of our significance by measurements made where the signal is expected to be indistinguishable from 0, we apply a matched filter using the NLA model shape $\mathbf{t}$. This weights bins by their expected signal:

$$S/N = \frac{\mathbf{t}^\top \mathbf{C}^{-1} \bar{\mathbf{d}}}{\sqrt{\mathbf{t}^\top \mathbf{C}^{-1} \mathbf{t}}}. \quad (16)$$

The statistic is independent of the empirical model amplitude and does not depend on the assumed cosmology. Equation 15 tests the data against a null of zero signal, and the matched filter of Equation 16 uses only the shape of the model, not its amplitude. To determine the significance over a given scale, we apply this only to the bins larger than $R_{\min}$. Combined over the three LRG redshift bins above $R > 200 h^{-1} \text{Mpc}$, the matched filter gives 3.7σ .

Our covariances are estimated from the scatter between sky regions, which in some cases are smaller than the measurement scales. Therefore, the very large-scale modes we measure span multiple regions, and we may underestimate large-scale covariance. However, we also expect the measurements to be close to shape-noise dominated, especially with the $\sigma_s = \frac{1}{\sqrt{2}}$ of multiplets, so this sample variance issue may not dramatically impact our detection significance.

To test this, we merge spatially adjacent regions into larger ones spanning the scales of interest and re-estimate the covariance. The origin of the merging grid is arbitrary and the resulting significance is moderately sensitive to it, so we enumerate several offsets and quote the median. We also test a range of region sizes. Across all samples, the number of merged regions ranges from the originally used 2500 in BGS2, down to tens of merged regions, corresponding to physical sizes of around $10 h^{-1} \text{Mpc}$ for the smallest BGS regions and up to $600 h^{-1} \text{Mpc}$ for ELG. Following the original region construction, each is built with the same number of multiplets and randoms are subtracted on a per-region basis. To see if any change in covariance reflects the inclusion of larger-scale modes rather than a lower number of regions, we repeat this but by randomly merging non-adjacent regions. These randomly-merged regions produce consistent covariance at all merging levels, as would be expected if the regions are independent, while adjacently-merged regions will depart from it at different merging levels. The difference between them isolates the contribution of modes larger than a single region.

For LRG3 and both ELG bins, whose regions are already comparable to the separations of interest, the adjacent and randomly merged results agree. Therefore, we keep the full number of regions that the measurements were originally made in for these samples. For the other samples, the effect appears once the regions are smaller than the separation being tested: for region sizes larger than $\sim 140 h^{-1} \text{Mpc}$, the merged and randomized results agree and below that the difference grows as region size shrinks. Since the combined LRG signal above $200 h^{-1} \text{Mpc}$ is dominated by LRG3, our large-scale detection significance for all LRGs is largely unaffected. For BGS, the significance above $200 h^{-1} \text{Mpc}$ falls from 4.1σ to 2.9σ . The randomized control does not reproduce this drop, indicating modes which supersede region scales. We also find that this significance does not converge as the number of regions becomes too small to estimate covariance. The lowest redshift BGS bins do not contain multiple independent regions which span more than $200 h^{-1} \text{Mpc}$. Therefore, we use the larger-region estimation and quote the BGS significances above $200 h^{-1} \text{Mpc}$ as an upper bound.

We note that these cannot be validated with the ABACUS-SUMMIT mocks, as the realizations involve replications of the box volume and are not suitable for characterizing the measurement scatter. A full analytic treatment, combining the Gaussian and super-sample covariance terms calibrated against the region-based estimate, would resolve this properly. However, this would require a density-response function for the multiplet alignment statistic that has not been derived; we leave this to future work.

Detection significances from the matched filter and estimated with the merged regions are shown on Figure 6, which displays the signals used to compute them. For LRGs, over scales of 100, 200, and $300 h^{-1} \text{Mpc}$, we find detection significances of 18.2σ , 3.7σ , and 1.8σ respectively. As mentioned above, the large-scale LRG signal is concentrated in the highest-redshift bin, 4.5σ above $200 h^{-1} \text{Mpc}$ when determined from LRG3 alone, while LRG1 and LRG2 are individually consistent with no large-scale signal (1.9σ , 0.3σ). We adopt the three-bin combination as the headline result, since the per-bin variation is consistent with cosmic variance.

Measurements for BGS and ELG were not made over $300 h^{-1} \text{Mpc}$ because we expect it would be an impractical use of computing resources. At these scales, the angular limits of the survey footprint for BGS and the low density of ELG multiplets will dilute significance well below that of LRGs. However, the high density of BGS (resulting in many more multiplets) produces more significant detections than LRGs below this. We detect tidal shear from BGS multiplets over 100 and $200 h^{-1} \text{Mpc}$ at $\leq 13.8\sigma$ and $\leq 2.9\sigma$. The equivalent for ELGs is 3.3σ and 0.4σ .

7.2. Comparison to Previous Measurements

Reliable measurements of intrinsic alignment on large scales are a crucial ingredient for many of the applications discussed in Section 1. Here we place our multiplet measurements in the context of the existing landscape of large-scale alignment detections.

Figure 9 compares the cumulative signal-to-noise of the strongest published alignment measurements that extend beyond projected separations of $100 h^{-1} \text{Mpc}$. These include the galaxy-shape intrinsic alignment of BOSS LOWZ (Singh et al. 2015, w_{g+}), the CMASS-UNIONS and eBOSS LRG-




FIG. 9.— The landscape of large-scale intrinsic tidal alignment detections. For each measurement we plot the cumulative signal-to-noise ratio accumulated from the largest measured scale inward. Solid, dotted, and dashed lines denote galaxy-shape intrinsic alignment measurements (w_{g+}); dash-dotted lines denote multiplet alignment ($\mathcal{E}_+$). All curves use the binning of the original published analyses. The cumulative statistic assumes diagonal covariance and therefore represents an upper bound on the true combined significance.

UNIONS measurements of Hervas Peters et al. (2025), and the DESI DR1 red BGS and high-redshift LRG measurements of Siegel et al. (2025). Against these we plot our DESI DR1 LRG multiplet result (Lamman et al. 2024b) and the DESI DR2 LRG, BGS, and ELG multiplet measurements presented in this work. We use the reported alignment signal and its uncertainty, in the published binning, for each measurement.

Several relevant measurements are not shown. Fourier-space analyses such as Kurita & Takada (2023) report a high total signal-to-noise for the E-mode intrinsic-alignment power spectrum, but that significance is accumulated across a broad range in wavenumber dominated by quasi-nonlinear scales, and a bandpower at a given k does not correspond to a single physical separation, so it cannot be placed on the horizontal axis of Figure 9. The BAO-scale intrinsic alignment of Xu et al. (2023) does reach large scales, but is measured in linear three-dimensional bins around the BAO feature rather than in projected separation, so it is not directly comparable; we note that its signal-to-noise is comparable to the CMASS-UNIONS measurement.

Although the galaxy-shape measurements report the projected correlation w_{g+} and the multiplet measurements report $\mathcal{E}_+$, the two can be compared on a signal-to-noise basis. The multiplet estimator averages a relative-orientation signal over multiplet-tracer pairs, while w_{g+} averages a shear over shape-density pairs; in both cases the uncertainty on the associated two-point clustering term is negligible relative to that on the alignment signal, so the per-bin signal-to-noise values are directly comparable even though the underlying estimators differ. The most significant difference between all measure-

ments is the choice of $\Pi_{\max}$. IA measurements are typically made with a flat $\Pi_{\max}$ across all bins of $60\text{--}100h^{-1}\text{Mpc}$. The multiplet measurements are the only ones using the suggested variable $\Pi_{\max}$ from Lamman et al. (2025), and therefore also represent an advancement in the alignment estimator.

Two caveats apply to this comparison. First, the per-bin signal-to-noise depends on the choice of binning; finer binning divides a fixed signal among more bins. The cumulative statistic in Figure 9 largely removes this sensitivity, since quadrature-summed signal-to-noise is approximately conserved under rebinning of independent bins. Second, and less tractably, neighboring bins are correlated, and the cumulative statistic neglects this off-diagonal covariance; it therefore overstates the true combined significance. Both effects should be borne in mind, but neither qualitatively changes the comparison.

The multiplet measurements reach a cumulative large-scale signal-to-noise more than an order of magnitude above the strongest galaxy-shape measurements over the same range of separations. We emphasize, however, that the galaxy-shape measurements shown here are drawn from earlier data sets; a like-for-like comparison against galaxy intrinsic alignment measured on DESI DR2 would be more equitable. We also note that an ELG-like galaxy population has never yielded a successful intrinsic-alignment detection, so the ELG multiplet measurement has no galaxy-shape counterpart to compare against.

Despite the challenges of comparing measurements made with different methods and data, Figure 9 demonstrates the remarkable growth of this field over the past decade. Improvements are driven by new estimators, the use of multiplet shapes, and the growing spectroscopic data available from large cosmological surveys.

7.3. Alignment Strength and Redshift

There is evidence, from both hydrodynamic simulations and observations, that the intrinsic alignment amplitude A_{IA} of individual galaxies increases with redshift over $0 < z < 1$ in LRG-like samples (Samuroff et al. 2021; Herle et al. 2026; Siegel et al. 2025). We might expect multiplet alignment to follow a similar trend, though it is not obvious that it should track the individual-galaxy result. Here we take an initial look at the redshift dependence of multiplet alignment, although a pure comparison across redshifts is limited by variations across and within samples, particularly the fact that our multiplet definitions vary between redshift bins in BGS. Figure 10 shows the measured amplitude of multiplet alignment as a function of redshift for all three tracers, reported both as the shear-response parameter τ and as the more common A_{IA} . For this comparison, the important distinction between the two is that A_{IA} includes the growth factor, $A_{\text{IA}} \propto D(z)$.

For BGS the interpretation is complicated by the sample itself. As described in Section 2, the flux-limited BGS sample changes substantially in both density and intrinsic luminosity across our five redshift bins, and we adopt a different multiplet definition in each bin to track the changing density (Section 3). The bias used to convert the measured signal to an amplitude is determined directly from the data (Section 5.4, Figure 8), and its uncertainty is propagated into the amplitude errors shown here. To assess the sensitivity of the amplitude to the multiplet definition, we additionally recompute it in each bin over a range of linking lengths between 1.5 and $8h^{-1}\text{Mpc}$ (Figure 1); the resulting spread is shown as the wider shaded




FIG. 10.— Redshift dependence of the multiplet alignment amplitude for the BGS, LRG, and ELG samples, shown as both the shear-response parameter τ and the conventional intrinsic-alignment amplitude A_{IA} . The two are related by a factor of the growth factor, $A_{IA} \propto \tau D(z)$ (Equation 10), so a sample whose τ is roughly constant with redshift can show a mild trend in A_{IA} . For BGS, the galaxy bias used to set the amplitude is measured directly from the data (Figure 8), and its uncertainty is included in the error bars. Because the BGS density varies strongly with redshift, the multiplet definition is also changed between BGS redshift bins; the wider shaded bars show the spread in amplitude obtained when the linking length is varied between 1.5 and $8 h^{-1} \text{Mpc}$ (Figure 1). Even so, the multiplets in each BGS bin differ intrinsically because of the changing density, so the BGS points should not be read as a clean evolution at fixed sample. The trend is more safely compared within the LRG and ELG samples, which hold a more nearly constant comoving density and a single linking length across their redshift bins. The star point marks the cross-correlation of ELG multiplets with LRG tracers over the overlapping range $0.8 < z < 1.1$, cleanly demonstrating that ELG multiplets are weaker tracers of intrinsic shear than LRGs.

error bars in Figure 10. Even with this accounted for, the multiplets in each BGS bin remain physically distinct because of the changing density, so the trend should not be read as a clean evolution at fixed sample. With that caveat, we do see the multiplet alignment amplitude increase with redshift across the BGS range, consistent with the trend reported for individual galaxies.

The redshift evolution is more safely compared within the LRG and ELG samples, which maintain a more nearly constant comoving density and use a single linking length across their redshift bins. Within LRGs, the direction of the apparent trend depends on which amplitude is considered: the multiplet A_{IA} decreases modestly with redshift while τ increases, reflecting the opposite growth-factor scaling of the two parameters. The overall picture is that the LRG multiplet alignment amplitude is fairly constant across this redshift range.

The ELG multiplets are similarly constant with redshift, but with a substantially lower amplitude than either BGS or LRG. This is notable. Individual ELGs show no measurable intrinsic alignment at all, but it was not a given that ELG multiplets would have a much lower amplitude than LRG multiplets, since the multiplet orientation does not depend on resolving individual galaxy shapes. We find that they do. A cleaner comparison is provided by the cross-correlation of ELG multiplets with LRG tracers over the overlapping redshift range $0.8 < z < 1.1$. The ELG multiplet amplitude remains much lower, demonstrating that the difference is not simply a matter of the tracer clustering, although the two multiplet definitions

still differ somewhat. This indicates that ELGs are found in less dense environments and are, perhaps, weaker tracers of the filamentary structure that sources multiplet alignment, consistent with expectations from morphology–environment correlations. This is a clear motivation to explore these trends in hydrodynamic simulations, where the astrophysical causes can be examined directly.

7.4. The BAO Feature

The signature of Baryon Acoustic Oscillations (BAO) in the tidal alignment signal can be understood through the NLA framework (Okumura et al. 2020; van Dompeler et al. 2023). A tracer at the center of the acoustic “ring” sees an excess of matter at the BAO scale distributed isotropically around it. This symmetric shell exerts no net tidal stretching in any preferred transverse direction, *diluting* radial alignment at separations near the sound horizon. At separations just inside and outside the shell, by contrast, the overdensity is anisotropic with respect to the tracer and *enhances* alignment. The net effect is a localized dip in alignment at the BAO scale relative to the smooth-model expectation. This has been measured in the galaxy density–ellipticity correlation by Xu et al. (2023), who reported $2 - 3\sigma$ evidence of the BAO feature using BOSS CMASS galaxies with DESI Legacy Imaging shape measurements. Here we assess whether an analogous feature is detectable in multiplet alignment.

We restrict this analysis to the BGS and LRG samples,




FIG. 11.— The average of 25 realizations of a mock DESI DR2 LRG catalog, compared to the NLA model predictions with and without an acoustic feature. Each model’s amplitude is fit to the measurements. The measurements are consistent with the BAO model, displaying a characteristic dip at the acoustic scale and peak beyond it.

in linear bins around the expected BAO feature, $55 < R < 180 h^{-1}\text{Mpc}$. The acoustic feature occupies a fixed comoving scale with an approximately constant width in R , and linear binning samples it more evenly than the scale-invariant logarithmic binning used earlier. To test for the presence of the feature, we compare two model templates for $\mathcal{E}_+(R)$, computed as in Section 4.2 but differing only in the input matter power spectrum: one using the full spectrum with acoustic oscillations, and one using a smooth “no-wiggle” spectrum with the oscillations removed. These are the same models used in DESI’s BAO main analysis (Andrade et al. 2025; Brieden et al. 2022). The amplitude of each model, τ , is fit to the measured signal independently and in each redshift bin.

We validated this model using the 25 mock realizations of the LRG sample, as described in Section 6. Their measurements, averaged over all realizations but kept in the three redshift bins, are compared to models in Figure 11. As discussed there, the amplitude of these mock measurements disagree with the true signal, but, as predicted by NLA, is consistent in its scale dependence. Here we clearly see the acoustic feature in all redshift bins, consistent with the model prediction.

For each redshift bin in the LRG and BGS sample, we form $\Delta\chi^2 = \chi_{\text{nw}}^2 - \chi_{\text{w}}^2$. A positive value indicates a preference for the acoustic template, zero indicates no acoustic information, and a negative value corresponds to a non-physical negative acoustic amplitude – which indicates a residual which is not described well by either template. We also fit the no-wiggle amplitude and apply a matched filter to the residual using the wiggle-minus-no-wiggle difference $\Delta m = m_{\text{w}} - m_{\text{nw}}$ as the template, yielding a signed detection significance $S/N = (\Delta m^T C^{-1} r) / \sqrt{\Delta m^T C^{-1} \Delta m}$, where r is the residual. We adopt this as our detection statistic, as it isolates the wiggle-specific information. Each redshift bin is treated independently, so their information combines at the covariance level: $C_{\text{comb}}^{-1} = \sum_s C_s^{-1}$. Because the acoustic feature is a small fraction of the total signal and lies in the most covariant bins, we validate the analytic significance against empirical null distributions built by bootstrap resampling and sign-flipping the sky regions; the empirical and analytic significances agree.




FIG. 12.— Tidal alignment of LRG multiplets around the acoustic scale. Each of the three redshift bins was independently used to fit the model amplitudes before being combined at the covariance level. These measurements display no preference for the model with a BAO feature over the broadband.




FIG. 13.— The correlation matrix corresponding to the measurements shown in Figure 12, and accounted for in our BAO analysis. Here the identity matrix has been subtracted for clarity. Adjacent bins are positively correlated, although not as strongly as the larger-scale measurements shown in Figure 7.

Figure 12 shows the inverse-variance-combined LRG signal against the two model templates, and Figure 13 the corresponding combined correlation matrix, with the identity subtracted to highlight the off-diagonal structure. We find no evidence for a BAO feature in the multiplet alignment signal. In the three redshift bins individually, the matched-filter significances are $S/N = -0.56, -0.64$, and -0.09 for the low, middle, and high-redshift LRG samples respectively; combined, $S/N = -0.78$. The data prefer the smooth template, with a combined $\Delta\chi^2$ of -9.2 , although the magnitude is well within the noise. Repeating the analysis with the merged regions of Section 7.1 reduces the sensitivity in the acoustic window by roughly 10% and shifts the combined matched-filter significance by less than the scatter between merging-grid choices, leaving the conclusion unchanged.

To see if this non-detection reflects the true absence of a feature rather than an insensitive measurement, we performed




FIG. 14.— Figure 12 but averaged over the BGS redshift bins. Although the measurements appear to have a feature similar to the BAO expectation, they systematically rise with R relative to both models. After fitting for an R -based correction, we do not find a consistent R dependence of the correction between redshift bins, indicating that this excess is not due to a single large angular systematic.

an injection-recovery test. We injected a synthetic BAO signal of amplitude A , in units of the model-predicted wiggle amplitude ($A = 1$ corresponding to the feature at the strength predicted for each sample’s fitted τ), coherently into the per-region measurement vectors, and re-ran the analysis. The results indicate that our measurements would be sensitive to a BAO feature at the model-predicted amplitude, recovering it at $\sim 3\sigma$. Its absence could indicate that the linear-alignment model is over-predicting the response of multiplet alignment to the BAO feature, although could simply be variance.

Figure 14 similarly shows the combined signal of all BGS redshift bins. In contrast to the LRG sample, the BGS measurements are not well described by either model template. A feature appears to be present here, but the signal rises above both the wiggle and no-wiggle predictions at large separations, and the broadband fit is only moderately good ($\chi^2/\text{dof} = 1.6$ for the wiggle model, 1.8 for no-wiggle, combined over 60 dof), with the weakest agreement concentrated in BGS3. This scale-dependent excess is strongest in the lowest-redshift bins. The combined measurements yield $S/N = -3.28$ and $\Delta\chi^2 = 13.1$, corresponding to 3.6σ . Although $\Delta\chi^2$ indicates a preference for the BAO feature, fitting amplitudes separately to each template makes the test biased by excess at large R , where the BAO feature appears as a peak over the smooth model.

This large-scale excess in the BGS signal could be a sign of unaccounted angular systematics, so we test for one by fitting an empirical, per-redshift-bin correction to absorb R -dependent systematics. For each redshift bin, we fit the wiggle and no-wiggle amplitudes as above, average the two τ -scaled model curves into a single broadband reference (so the small acoustic wiggle has negligible leverage), and fit a linear correction function to the data’s residual against that reference. This is fit by generalized least squares under the same region-block covariance used throughout, and subtracted from the data on top of the existing per-region randoms subtraction. For this test we extend the R bins to $55 < R < 210 h^{-1}\text{Mpc}$ to capture the excess where it is largest.

The correction improves the fit for both models, but neither becomes an acceptable description of the data. Notably, the fitted corrections are not consistent across redshifts. BGS1-

3 require a growing positive term at large R while BGS4-5 require a negative one. This argues against a single large angular systematic. Some residual geometric contamination cannot be ruled out, but the excess is equally consistent with noise.

We conclude that a robust acoustic measurement in the alignment signal for both LRG and BGS is better pursued with 3D estimators and the greater statistical power of a future DESI data release.

8. CONCLUSION

We present the largest-scale signal of intrinsic tidal shear yet detected. Multiplets of galaxies in DESI’s DR2 LRG samples display significant alignment with the surrounding matter field at separations over $200 h^{-1}\text{Mpc}$, and we find alignment of BGS and ELG multiplets over $100 h^{-1}\text{Mpc}$. These represent an order of magnitude improvement over previous alignment measurements, reflecting advancements in data, measurement methods, and new multiplet definitions.

We also present a sample-dependent method for identifying multiplets, which increases the maximum separation between multiplet members as galaxy density decreases. This results in larger multiplets relative to previous work and a higher signal-to-noise ratio of alignment measurements. We find no improvement from quantifying multiplets with a full ellipticity over a single orientation.

Although a weaker signal, the ELG measurements are especially notable as they present the opportunity to make alignment measurements in samples previously excluded from IA applications. They are also the sample with the most improved measurements over DR1, mostly due to the new multiplet definition, and the sample with the most potential for improvement as ELG density increases with future surveys.

We additionally measure, for the first time, the multiplet autocorrelation. Results are not consistent with model expectations calibrated to the cross-correlations, likely due to an underestimate of the galaxy bias. However, the autocorrelation provides the opportunity to estimate the linear galaxy bias directly from our samples, which we do for BGS.

Taking into account the uncertainty in galaxy bias and variations from a changing multiplet definition, we find the alignment strength of BGS multiplets to be positively correlated with redshift. Less significant redshift correlations are found in LRGs and ELGs, which may be related to these samples having a more consistent comoving density. In their redshift overlap, we compare the alignment strength of LRG and ELG multiplets by measuring them both relative to the full LRG sample. We find that ELGs have half the alignment amplitude of LRGs, demonstrating that the small-scale clustering of spiral galaxies is less correlated with their large-scale environment than that of large elliptical galaxies. It would be interesting to further examine these trends with hydrodynamic simulations.

Finally, we examine the potential presence of a BAO feature in LRG and BGS. The acoustic dip is present in mock data and an injection-recovery test indicates that our measurements should be sensitive to it. However we find no detection, even after a first-order accounting of an unknown angular systematic. This could still be a result of angular effects unaccounted for in the survey mask, or merely sample variance. We leave a more detailed analysis to future samples and a more suitable (3D) estimator.

Multiplets have now enabled the largest-scale detections of

intrinsic tidal shear, and their orientations are robust to the imaging systematics that plague direct detections of IA with individual galaxy shapes. This unlocks the potential of intrinsic alignments as a practical cosmological probe. Especially promising applications involve exploring primordial features that have since inflated to the very large scales that multiplet alignment is sensitive to.

ACKNOWLEDGMENTS

CL acknowledges useful discussions with Christopher Hirata and the CCAPP Cosmo Group. CL is supported by an NSF Postdoctoral Fellowship under award 2502789.

This material is based upon work supported by the U.S. Department of Energy (DOE), Office of Science, Office of High-Energy Physics, under Contract No. DE-AC02-05CH11231, and by the National Energy Research Scientific Computing Center, a DOE Office of Science User Facility under the same contract. Additional support for DESI was provided by the U.S. National Science Foundation (NSF), Division of Astronomical Sciences under Contract No. AST-0950945 to the NSF's National Optical-Infrared Astronomy Research Laboratory; the Science and Technology Facilities Council of the United Kingdom; the Gordon and Betty Moore Foundation; the Heising-Simons Foundation; the French Alternative Energies and Atomic Energy Commission (CEA); the Secretariat of Science, Humanities, Technology and Innovation (SECIHTI) of Mexico; the Ministry of Science, Innovation and Universities of Spain (MICIU/AEI/10.13039/501100011033), and by the DESI Member Institutions: <https://www.desi.lbl>.

gov/collaborating-institutions. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the U. S. National Science Foundation, the U. S. Department of Energy, or any of the listed funding agencies.

The authors are honored to be permitted to conduct scientific research on I'oligam Du'ag (Kitt Peak), a mountain with particular significance to the Tohono O'odham Nation.

Generative AI (Anthropic's Claude) was used in this work to assist with software development: optimizing existing functions, generating batch scripts, editing plotting scripts, writing unit tests, debugging, and testing pipelines for modeling and significance estimation. Final measurement code that included AI-generated changes was validated against the original pipeline and an independent TreeCorr-based implementation. The modeling and significance code was tested for consistency with previously validated human-made code and analytic computations.

DATA AVAILABILITY

[Zenodo will appear here for final reading]

Code used in this work is publicly available in the repository github.com/cmlamman/spec-IA.

REFERENCES

- Achúcarro A., et al., 2022, *Inflation: Theory and Observations*, doi:10.48550/arXiv.2203.08128, <https://ui.adsabs.harvard.edu/abs/2022arXiv220308128A>
- Akitsu K., Kurita T., Nishimichi T., Takada M., Tanaka S., 2021, *Physical Review D*, 103, 083508
- Andrade U., et al., 2025, *Physical Review D*, 112, 083512
- Bartelmann M., Schneider P., 2001, *Physics Reports*, 340, 291
- Bridle S., King L., 2007, *New Journal of Physics*, 9, 444
- Brieden S., Gil-Marín H., Verde L., 2022, *Journal of Cosmology and Astroparticle Physics*, 2022, 024
- Brown M. L., Taylor A. N., Hambly N. C., Dye S., 2002, *Monthly Notices of the Royal Astronomical Society*, 333, 501
- Chisari N. E., 2025, *Astronomy and Astrophysics Review*, 33, 5
- Chisari N. E., Dvorkin C., 2013, *Journal of Cosmology and Astroparticle Physics*, 12, 029
- DESI Collaboration et al., 2016, arXiv e-prints, p. arXiv:1611.00037
- DESI Collaboration et al., 2022, *The Astronomical Journal*, 164, 207
- DESI Collaboration ., et al., 2025, *Physical Review D*, 112, 083515
- DESI Collaboration et al., 2026, *J* 10.3847/1538-3881/ae4c43, 171, 285
- Dey A., et al., 2019, *The Astronomical Journal*, 157, 168
- Fernández-García E., et al., 2026, *Journal of Cosmology and Astroparticle Physics*, 2026, 002
- Guy J., et al., 2023, *The Astronomical Journal*, 165, 144
- Hahn C., et al., 2023, *The Astronomical Journal*, 165, 253
- Hartlap J., Simon P., Schneider P., 2007, *Astronomy and Astrophysics*, 464, 399
- Herle A., Chisari N. E., Hoekstra H., Neumann D., 2026, *Assembly bias and the redshift evolution of intrinsic alignments for LRGs*, doi:10.48550/arXiv.2607.00785, <https://ui.adsabs.harvard.edu/abs/2026arXiv260700785H>
- Hervas Peters F., et al., 2025, *Astronomy and Astrophysics*, 699, A201
- Ichikawa M., Shi J., Kurita T., Takada M., Nishimichi T., Blot L., 2026, arXiv e-prints, p. arXiv:2607.11114
- Jarvis M., Bernstein G., Jain B., 2004, *Monthly Notices of the Royal Astronomical Society*, 352, 338
- Joachimi B., Mandelbaum R., Abdalla F. B., Bridle S. L., 2011, *Astronomy and Astrophysics*, 527, A26
- Joachimi B., et al., 2015, *Space Science Reviews*, 193, 1
- Johnston H., et al., 2019, *Astronomy and Astrophysics*, 624, A30
- Kumwembe A. A., et al., 2025, *Research Notes of the American Astronomical Society*, 9, 305
- Kurita T., Takada M., 2023, *Physical Review D*, 108, 083533
- Kurita T., Jamieson D., Komatsu E., Schmidt F., 2025, *Parity Violation in Galaxy Shapes: Primordial Non-Gaussianity*, doi:10.48550/arXiv.2509.08787, <https://ui.adsabs.harvard.edu/abs/2025arXiv250908787K>
- Lamman C., Tsaprazi E., Shi J., Šarčević N. N., Pyne S., Legnani E., Ferreira T., 2024a, *The Open Journal of Astrophysics*, 7
- Lamman C., et al., 2024b, *Monthly Notices of the Royal Astronomical Society*, 534, 3540
- Lamman C., Blazek J., Eisenstein D. J., 2025, *Optimal intrinsic alignment estimators in the presence of redshift-space distortions*, doi:10.48550/arXiv.2504.16076, <https://ui.adsabs.harvard.edu/abs/2025arXiv250416076L>
- Maksimova N. A., Garrison L. H., Eisenstein D. J., Hadzhiyska B., Bose S., Satterthwaite T. P., 2021, *Monthly Notices of the Royal Astronomical Society*, 508, 4017
- Mandelbaum R., Hirata C. M., Ishak M., Seljak U., Brinkmann J., 2006, *Monthly Notices of the Royal Astronomical Society*, 367, 611
- Mena-Fernández J., et al., 2025, *Journal of Cosmology and Astroparticle Physics*, 2025, 133
- Miller T. N., et al., 2024, *The Astronomical Journal*, 168, 95
- More S., Kravtsov A. V., Dalal N., Gottlöber S., 2011, *The Astrophysical Journal Supplement Series*, 195, 4
- Myers A. D., et al., 2023, *The Astronomical Journal*, 165, 50
- Okumura T., Taruya A., Nishimichi T., 2020, *Monthly Notices of the Royal Astronomical Society*, 494, 694
- Paul P., Clarkson C., Maartens R., 2024, *Physical Review Letters*, 133, 121001
- Poppett C., et al., 2024, *The Astronomical Journal*, 168, 245
- Raichoor A., et al., 2023, *The Astronomical Journal*, 165, 126
- Ross A. J., et al., 2025, *Journal of Cosmology and Astroparticle Physics*, 2025, 125
- Samuroff S., Mandelbaum R., Blazek J., 2021, *Monthly Notices of the Royal Astronomical Society*, 508, 637

Samuroff S., et al., 2023, *Monthly Notices of the Royal Astronomical Society*, 524, 2195
 Schlafly E. F., et al., 2023, *The Astronomical Journal*, 166, 259
 Siegel J., et al., 2025, arXiv e-prints, p. arXiv:2507.11530
 Silber J. H., et al., 2023, *The Astronomical Journal*, 165, 9
 Singh S., Mandelbaum R., 2016, *Monthly Notices of the Royal Astronomical Society*, 457, 2301
 Singh S., Mandelbaum R., More S., 2015, *Monthly Notices of the Royal Astronomical Society*, 450, 2195

The LSST Dark Energy Science Collaboration et al., 2018, *The LSST Dark Energy Science Collaboration (DESC) Science Requirements Document*, doi:10.48550/arXiv.1809.01669,
<https://ui.adsabs.harvard.edu/abs/2018arXiv180901669T>
 Troxel M. A., Ishak M., 2015, *Physics Reports*, 558, 1
 Vaisakh R., et al., 2026, *DESI DR2 Reference Mocks: Clustering results from UCHUU ELGs and QSOs*, doi:10.48550/arXiv.2606.28559,
<https://ui.adsabs.harvard.edu/abs/2026arXiv260628559V>
 Xu K., Jing Y. P., Zhao G.-B., Cuesta A. J., 2023, *Nature Astronomy*, 7, 1259
 Zhou R., et al., 2023, *The Astronomical Journal*, 165, 58
 van Dompseleer D., Georgiou C., Chisari N. E., 2023, *The Open Journal of Astrophysics*, 6, 19

AUTHOR AFFILIATIONS

¹Center for Cosmology and AstroParticle Physics (CCAPP), Ohio State University, Columbus, OH 43210
²Department of Astronomy, The Ohio State University, Columbus, OH 43210, USA
³Department of Physics, The Ohio State University, Columbus, OH 43210, USA
⁴The Ohio State University, Columbus, 43210 OH, USA
⁵Lawrence Berkeley National Laboratory, 1 Cyclotron Road, Berkeley, CA 94720, USA
⁶Department of Physics, Boston University, 590 Commonwealth Avenue, Boston, MA 02215 USA
⁷Instituto Avanzado de Cosmología A. C., San Marcos 11 - Atenas 202. Magdalena Contreras. Ciudad de México C. P. 10720, México
⁸Instituto de Ciencias Físicas, Universidad Nacional Autónoma de México, Av. Universidad s/n, Cuernavaca, Morelos, C. P. 62210, México
⁹Department of Physics & Astronomy, University College London, Gower Street, London, WC1E 6BT, UK
¹⁰Departamento de Astrofísica, Universidad de La Laguna (ULL), E-38206, La Laguna, Tenerife, Spain
¹¹Instituto de Astrofísica de Canarias, C/ Vía Láctea, s/n, E-38205 La Laguna, Tenerife, Spain
¹²Instituto de Física, Universidad Nacional Autónoma de México, Circuito de la Investigación Científica, Ciudad Universitaria, Cd. de México C. P. 04510, México
¹³Center for Astrophysics | Harvard & Smithsonian, 60 Garden Street, Cambridge, MA 02138, USA
¹⁴Institut d'Estudis Espacials de Catalunya (IEEC), c/ Esteve Terradas 1, Edifici RDIT, Campus PMT-UPC, 08860 Castelldefels, Spain
¹⁵Institute of Cosmology and Gravitation, University of Portsmouth, Dennis Sciana Building, Portsmouth, PO1 3FX, UK
¹⁶Institute of Space Sciences, ICE-CSIC, Campus UAB, Carrer de Can Magrans s/n, 08913 Bellaterra, Barcelona, Spain
¹⁷Fermi National Accelerator Laboratory, PO Box 500, Batavia, IL 60510, USA
¹⁸Department of Physics, The University of Texas at Dallas, 800 W. Campbell Rd., Richardson, TX 75080, USA
¹⁹Department of Astronomy & Astrophysics, University of Toronto, Toronto, ON M5S 3H4, Canada
²⁰Department of Physics and Astronomy, University of California, Irvine, 92697, USA
²¹Departament de Física, Serra Hünter, Universitat Autònoma de Barcelona,

08193 Bellaterra (Barcelona), Spain
²²Institut de Física d'Altes Energies (IFAE), The Barcelona Institute of Science and Technology, Edifici Cn, Campus UAB, 08193, Bellaterra (Barcelona), Spain
²³NSF NOIRLab, 950 N. Cherry Ave., Tucson, AZ 85719, USA
²⁴Institució Catalana de Recerca i Estudis Avançats, Passeig de Lluís Companys, 23, 08010 Barcelona, Spain
²⁵Department of Physics and Astronomy, University of Waterloo, 200 University Ave W, Waterloo, ON N2L 3G1, Canada
²⁶Perimeter Institute for Theoretical Physics, 31 Caroline St. North, Waterloo, ON N2L 2Y5, Canada
²⁷Waterloo Centre for Astrophysics, University of Waterloo, 200 University Ave W, Waterloo, ON N2L 3G1, Canada
²⁸Instituto de Astrofísica de Andalucía (CSIC), Glorieta de la Astronomía, s/n, E-18008 Granada, Spain
²⁹Departament de Física, EEBE, Universitat Politècnica de Catalunya, c/Eduard Maristany 10, 08930 Barcelona, Spain
³⁰Université Clermont-Auvergne, CNRS, LPCA, 63000 Clermont-Ferrand, France
³¹Department of Physics and Astronomy, Sejong University, 209 Neungdong-ro, Gwangjin-gu, Seoul 05006, Republic of Korea
³²Queensland University of Technology, School of Chemistry & Physics, George St, Brisbane 4001, Australia
³³CIEMAT, Avenida Complutense 40, E-28040 Madrid, Spain
³⁴Max Planck Institute for Extraterrestrial Physics, Gießenbachstraße 1, 85748 Garching, Germany
³⁵University of Michigan, 500 S. State Street, Ann Arbor, MI 48109, USA
³⁶Departamento de Física, Universidad de los Andes, Cra. 1 No. 18A-10, Edificio Ip, CP 111711, Bogotá, Colombia
³⁷Observatorio Astronómico, Universidad de los Andes, Cra. 1 No. 18A-10, Edificio H, CP 111711 Bogotá, Colombia

This paper was built using the Open Journal of Astrophysics L^AT_EX template. The OJA is a journal which provides fast and easy peer review for new papers in the astro-ph section of the arXiv, making the reviewing process simpler for authors and referees alike. Learn more at <http://astro.theoj.org>.